


Parallelizing Mandelbrot

Overview

For improving the performance of our mandelbrot set calculations, we used a variety of strategies, including

- Adjusting computation distribution
 - Per-core distribution of workload
 - (OpenMP or CUDA)
 - between CPU and GPU
 - (let CPU contribute with idle time as GPU calculation progresses)
- Caching results
 - Useful when moving the viewing window, especially with horizontal/vertical shifts


Group 8:
Patrick Leiser,
Jared Pugh,
Catherine Yaroslavtseva



Distribution


- Sample code - block distribution

- Unbalanced workload
 - Depending on the iterations required on a given row, some take longer to calculate
- Balance with dynamically-scheduled distribution
 - Accomplished through dynamic scheduling in OpenMP
 - Because number of iterations for a given coordinate is hard to predict, a cyclic distribution doesn't help much
 - Thus dynamic scheduling performs much better than either cyclic or block distributions




CPU vs. GPU Balancing

- Timesave from GPU is nonexistent with small problem size
 - As problem size increases, GPU starts to become more useful
 - A hybrid approach using both GPU and CPU computing can be faster than either alone
 - Calibrate a balance between GPU and CPU workload
 - Shift to GPU as workload increases
 - CPU also able to get bigger time-saves from caching
 - How colors are stored
 - And more details on next slide...



Cache

- Default implementation recalculates colors for every transformation
 - Results in redundant calculations
 - x/y values that have already been solved
 - Color values for previously calculated iterations
- Two potential cache implementations
 - Map colors to pixels indices
 - Results in significant time savings after the initial render
 - $O(1)$ cache access time with hashmap
 - Map colors to iteration values
 - May help especially with frequently encountered colors
 - Limited speedup because we use a simple (fast) color calculation
 - Just multiplication and division, no sine function like some examples use



Next Steps

- Process constant amounts of iterations and cache the results to produce a progressive drawing
- Dynamically choose distribution between CPU and GPU
 - Have them scan from opposite directions, meet in the middle
- Enable caching for the reflection and zoom features
- Preemptively pre-caching nearby out-of-frame values
- Utilize multiple GPUs?
- Combine all our optimizations

